

Hashed Wheel Timer

1. 简介
2. 复杂度分析
3. 数据结构
4. 实现细节
5. 性能
6. Hashed Hierarchical Timing Wheel

1. 简介

Hashed Wheel Timer 主要用来高效处理大量定时任务， 可以对任务进行高效的schedule和unschedule操作。

在Netty, Kafka中都有应用。

- [Kafka Timer](#)
- [Netty Timer](#)

2. 复杂度分析

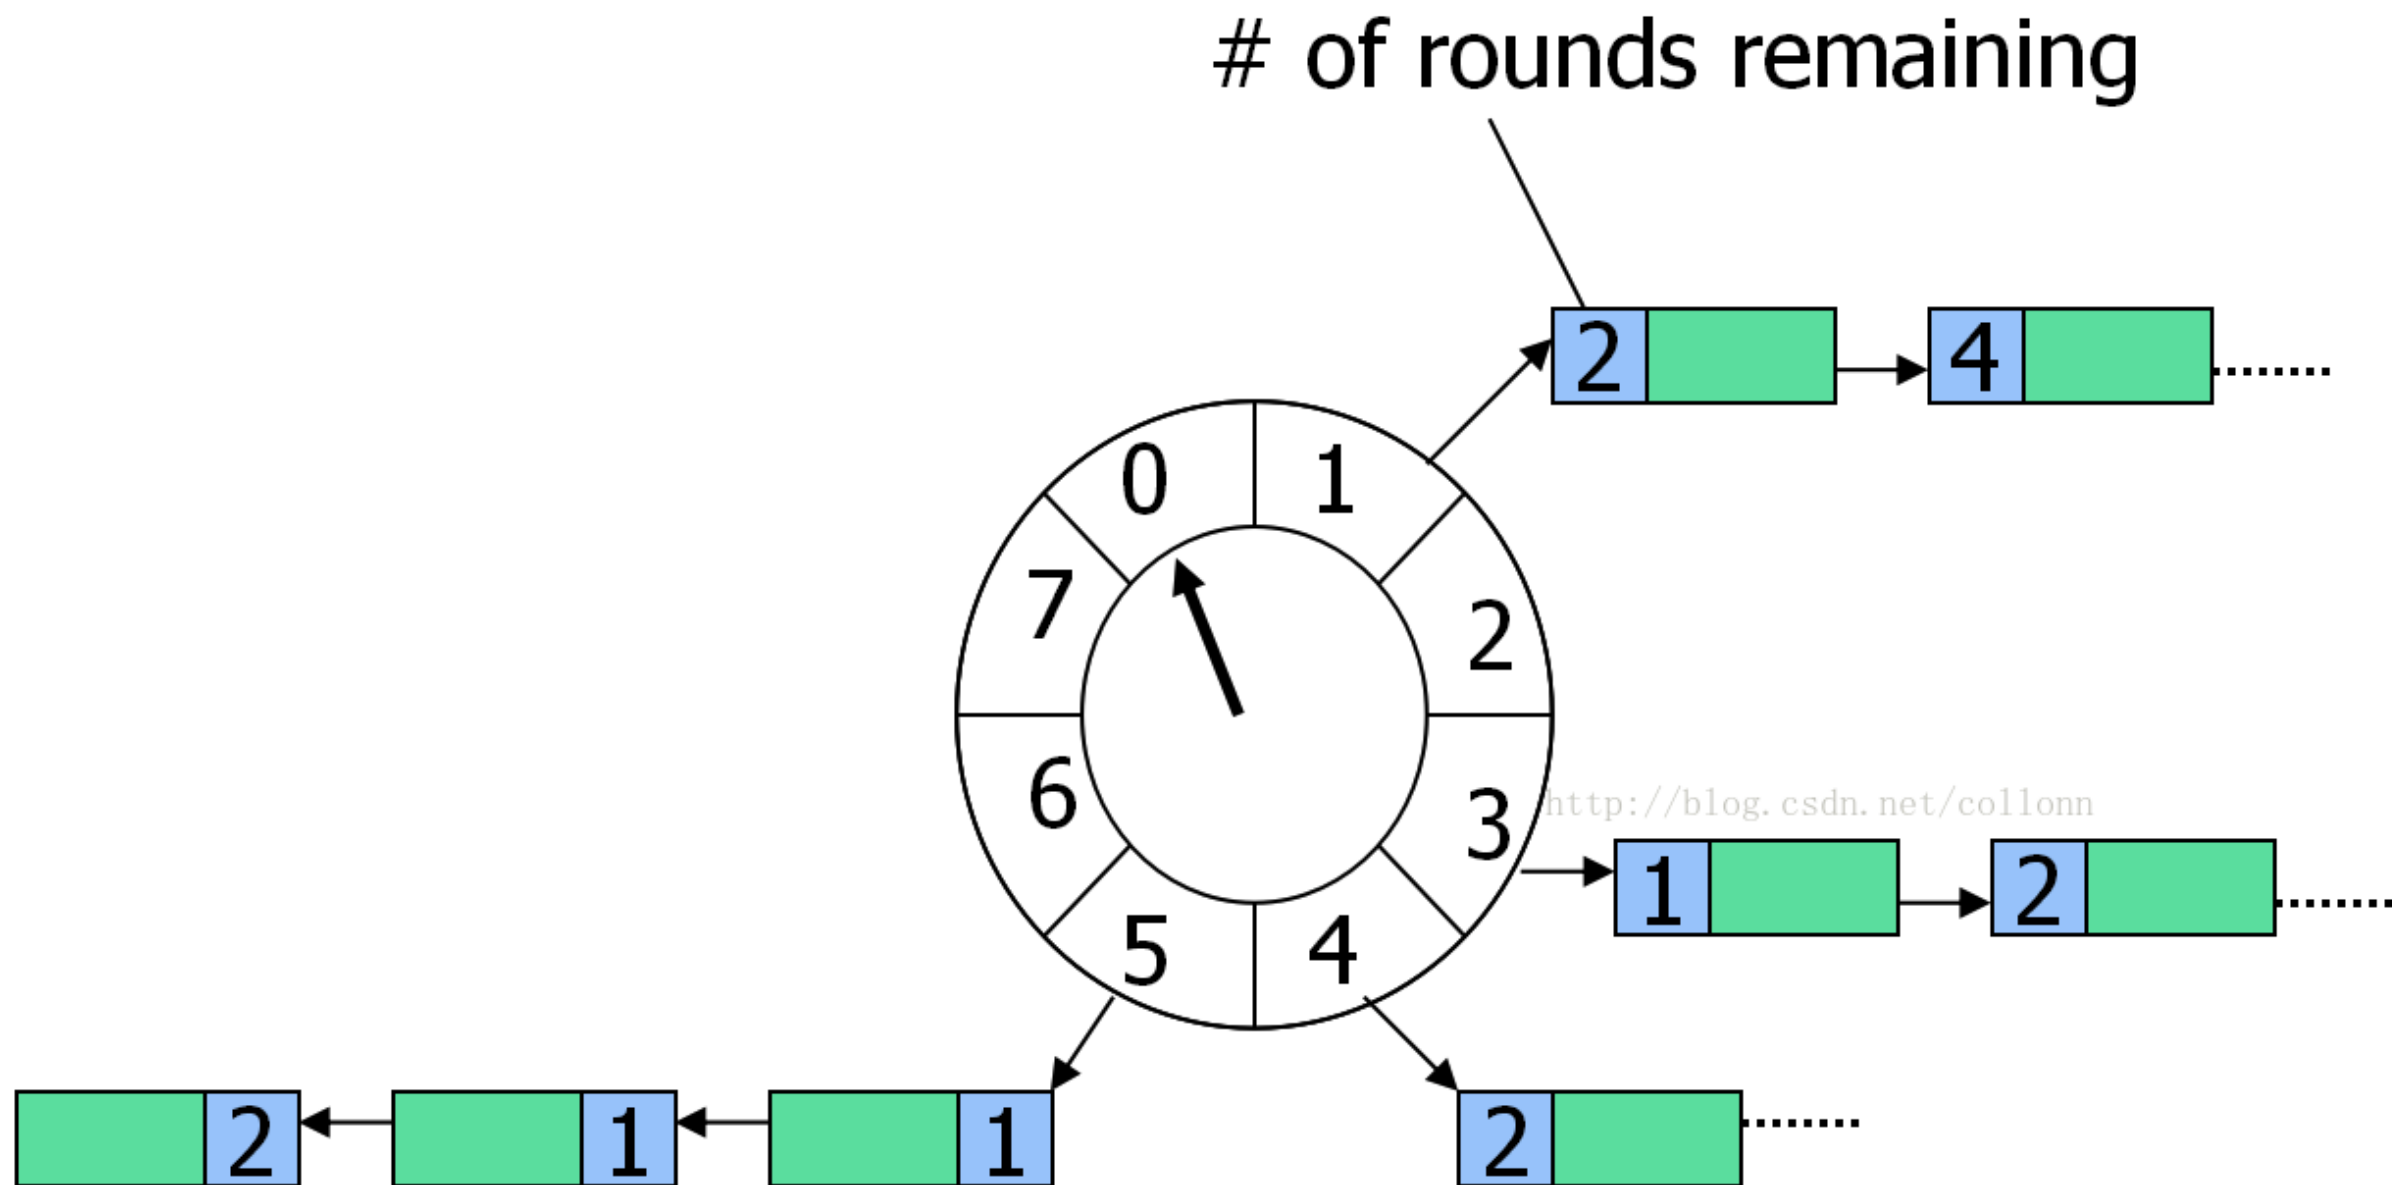
ScheduledThreadPoolExecutor

- $O(\log N)$ for adding new task
- $O(1)$ per each timer tick
- $O(\log N)$ cancelling the task

HashedWheelTimer

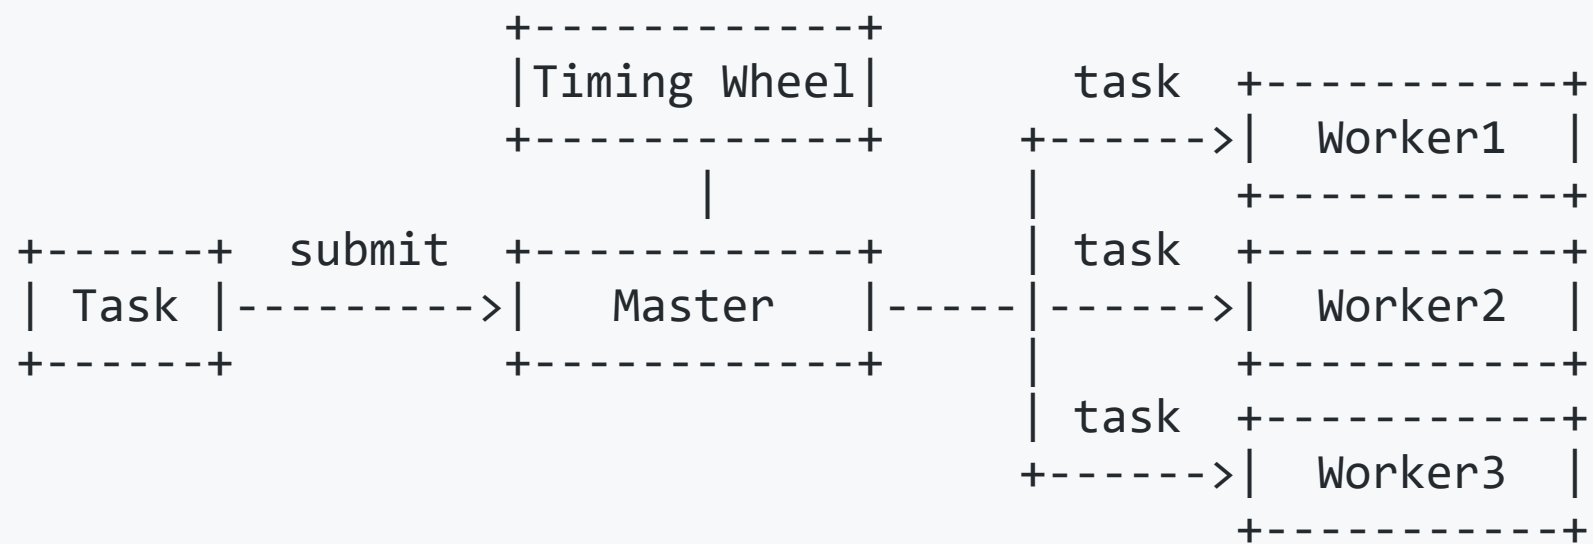
- $O(1)$ for adding new task
- $O(m)$ per each timer tick ($m = \text{bucket.size}$)
- $O(1)$ for cancelling a task

3. 数据结构



4. 实现细节

4.1. 线程模型



4.2. 实现细节

4.2.1. Timing wheel

```
while (running.get()) {  
    Set<Task<?>> bucket = ring[cursor];  
    for (Task<?> task : bucket) {  
        if (task.isCancelled()) {  
            bucket.remove(task);  
        } else if (task.rounds == 0) {  
            execute(task); // submit to worker thread  
            bucket.remove(task);  
            if (task.isReschedule()) // scheduleAtFixedRate  
                task.reschedule();  
        } else {  
            task.rounds--  
        }  
    }  
    this.idle.delay(tick);  
    cursor = (cursor + 1) & (size - 1)  
}
```

4.2.2. Task

```
public interface Task<T> extends ScheduledFuture<T>, Runnable {  
    boolean isReschedule(); void reschedule();  
}
```

4.2.3. schedule

```
public <T> ScheduledFuture<T> schedule(Callable<T> callable, long delay, TimeUnit u) {  
    final int offset1 = (int) (delay / tick), rounds1 = (offset1 / size);  
    Task<T> task = new OneShotTask<T>(rounds1, offset1, callable)  
    ring[(cursor+offset1) & (size-1)].add(task);  
    return task;  
}
```

```
public boolean isReschedule() {  
    return false;  
}  
  
public void reschedule() {  
    throw new UnsupportedOperationException();  
}
```


4.2.4. scheduleAtFixedRate

```
public ScheduledFuture<?> scheduleAtFixedRate(Runnable r, long d, long p, TimeUnit u){
    final int offset1 = (int) (d / tick), rounds1 = (offset1 / size);
    final int offset2 = (int) (p / tick), rounds2 = (offset2 / size);
    Task<T> task = new FixedRateTask(rounds1, offset1, r, rounds2, offset2);
    ring[(cursor+offset1) & (size-1)].add(task);
    return task;
}
```

```
public boolean isReschedule() {
    return true;
}

public void reschedule() {
    this.canceled.set(false);
    this.rounds1 = this.rounds2;
    this.offset1 = this.offset2;
    ring[(cursor+offset1) & (size-1)].add(task);
}
```

4.2.5. scheduleWithFixedDelay

```
public ScheduledFuture<?> scheduleWithFixedDelay(Runnable r, long d, long p, TimeUnit u)
    final int offset1 = (int) (d / tick), rounds1 = (offset1 / size);
    final int offset2 = (int) (p / tick), rounds2 = (offset2 / size);
    Task<T> task = new FixedDelayTask(rounds1, offset1, r, rounds2, offset2);
    ring[(cursor+offset1) & (size-1)].add(task);
    return task;
}
```

```
public boolean isReschedule() {
    return false; // see #run()
}

public void run() {
    super.run(); reschedule();
}

public void reschedule() {
    this.canceled.set(false);
    this.rounds1 = this.rounds2;
    this.offset1 = this.offset2;
    ring[(cursor+offset1) & (size-1)].add(task);
}
```

5. 性能

Benchmark	Mode	Cnt	Score	Error	Units
benchmarkScheduledExecutor	thrpt	200	795377.428 ± 176042.240		ops/s
benchmarkXScheduledExecutor (100ms tick, 512 wheel size)	thrpt	200	7452524.500 ± 79330.412		ops/s

6. Hashed Hierarchical Timing Wheel

